

Neural Preset for Color Style Transfer

(CVPR 2023)

Paper Authors:

Zhanghan Ke, Yuhao Liu, Lei Zhu, Nanxuan Zhao, Rynson W.H. Lau

Presentation author:

Umut Özyurt



MIDDLE EAST TECHNICAL UNIVERSITY

This presentation was prepared for the CENG796 - Deep Generative Models course for the Spring 2024 term.
It is part of the coursework assigned by our lecturer, Assoc. Prof. Ramazan Gokberk Cinbis.



Input Image (8K Resolution)



Ours (0.061s)


Output examples from the presented paper

Figures (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU



Output examples from the presented paper

Overview of the presentation

- Introduction
- Related work
- Method
- Experiments
- Main Contribution and Pros/Cons
- Conclusion
- Q & A



1- Introduction



Problem Definition

- Transfer only colors from an image without altering the targets texture. Hence, it is NOT artistic style transfer.
- Utilize self-supervised learning.
- Make the inference time fast.
- Use feasible GPU memory.




Figure 1. Example color style transfers to demonstrate how well the texture remained as it should.

Figure 1 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU


Why Color Style Transfer?

- Producing different styles of images, just like using different customizable Instagram filters.
- Further utilizing the 2-stage transfer approach, switching styles are genuinely fast.
- Using the models on different tasks, such as low-light image enhancement , underwater image correction , image dehazing, and image harmonization.



Application Examples

Low-Light Image Enhancement



Underwater Image Correction



Input Image

Reference

Ours


Figure 2. Applications of color style transfer model on Low-Light Image Enhancement and Underwater Image Correction Tasks

Application Examples

Image Dehazing



Image Harmonization



Input Image

Reference

Ours

Figure 3. Applications of color style transfer model on Image Dehazing and Image Harmonization tasks.

2- Related Work



Related Work

- Color Style Transfer
- Deterministic Color Mapping with CNNs.
- Self-Supervised Learning (SSL).



Related Work Discussion

- What is missing or should be improved so far?
- Visual artifacts
- High memory requirements, high inference times.
- Inefficient style switch for images and low scalability for high resolution.



3- Method



Two-Stage Pipeline




Figure 4. Proposed two-stage pipeline from the paper

Two-Stage Pipeline

(a) Color Normalization Stage




Figure 4. Proposed two-stage pipeline from the paper

Two-Stage Pipeline

(a) Color Normalization Stage

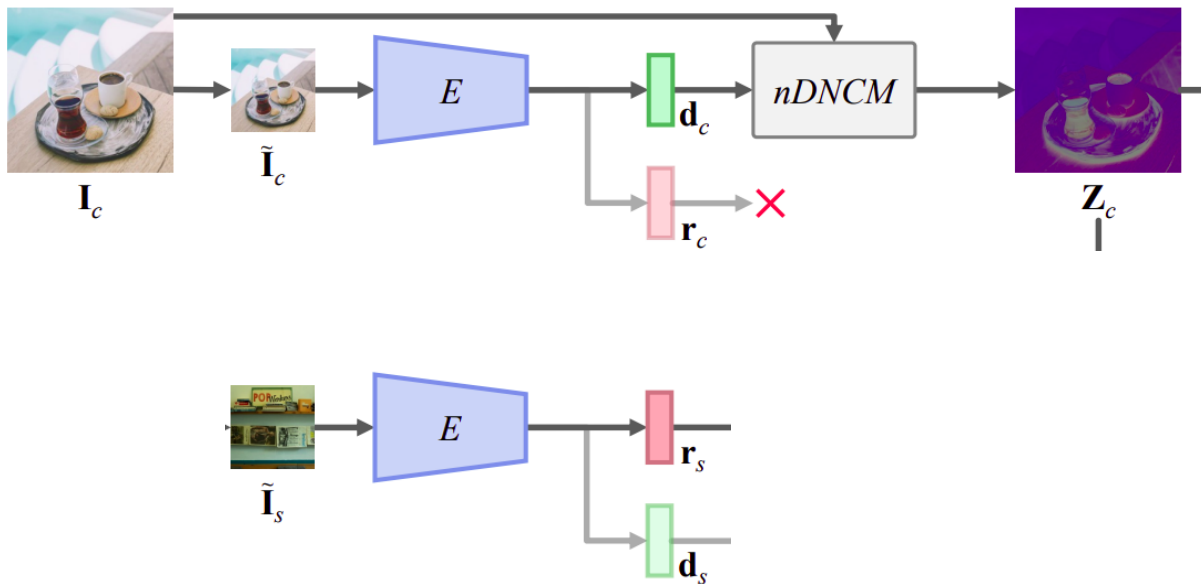



Figure 4. Proposed two-stage pipeline from the paper

Two-Stage Pipeline

(a) Color Normalization Stage



(b) Color Stylization Stage






Figure 4. Proposed two-stage pipeline from the paper

Two-Stage Pipeline

(a) Color Normalization Stage



(b) Color Stylization Stage



(c) Applying Color Style Presets




Figure 4. Proposed two-stage pipeline from the paper

Figure 4 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

DNCM Mapping

- What is DNCM?
- Deterministic Neural Color Mapping (DNCM) is proposed in the paper to model arbitrary deterministic color mappings
- It simply uses 3 matrix multiplications.



DNCM Mapping Visualization

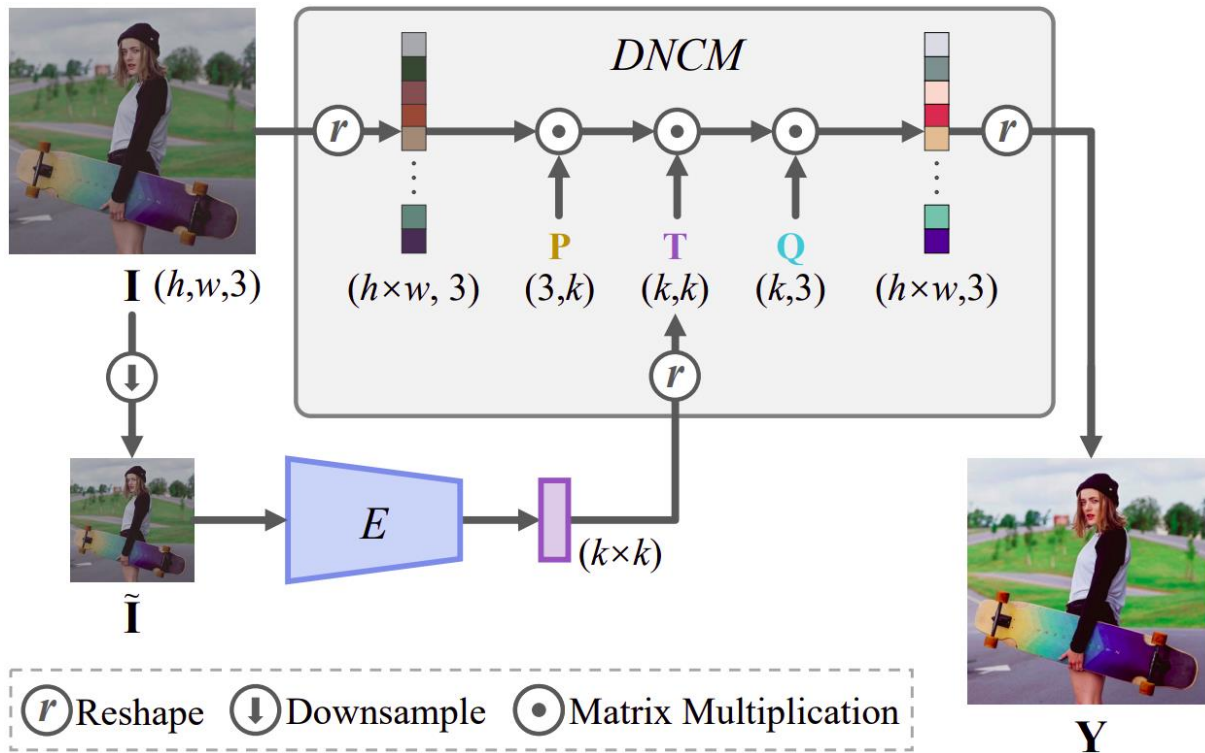


Figure 5. Proposed DNCM Module

Figure 5 (from the presented paper):


https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

Full Architecture

(a) Color Normalization Stage



(b) Color Stylization Stage




Figure 4. Proposed two-stage pipeline from the paper

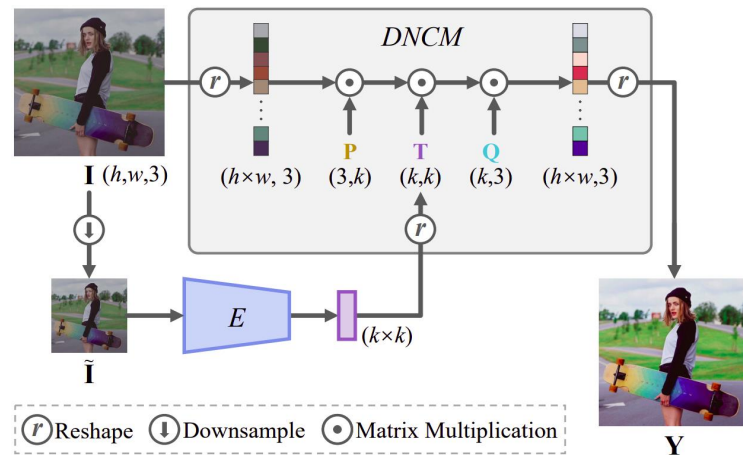
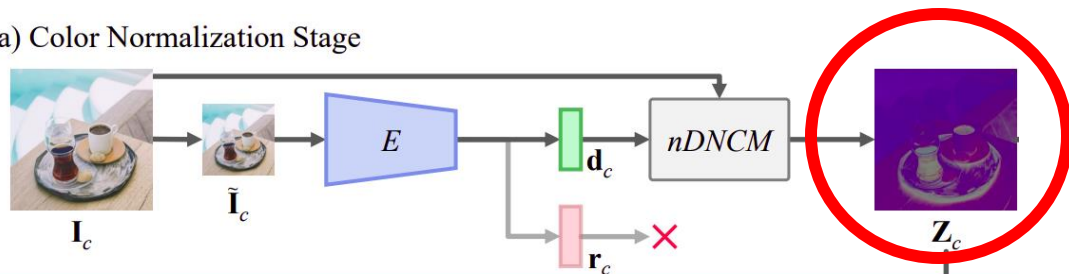


Figure 5. Proposed DNCM Module

Full Architecture

(a) Color Normalization Stage



Reusable outputs!

Combining any Z_c with any r_s is possible!

(b) Color Stylization Stage




Figure 4. Proposed two-stage pipeline from the paper

Figure 4, 5 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf




METU

How to Train the Models?

→ What parameters do we have to train?

(a) Color Normalization Stage



(b) Color Stylization Stage




Figure 4. Proposed two-stage pipeline from the paper

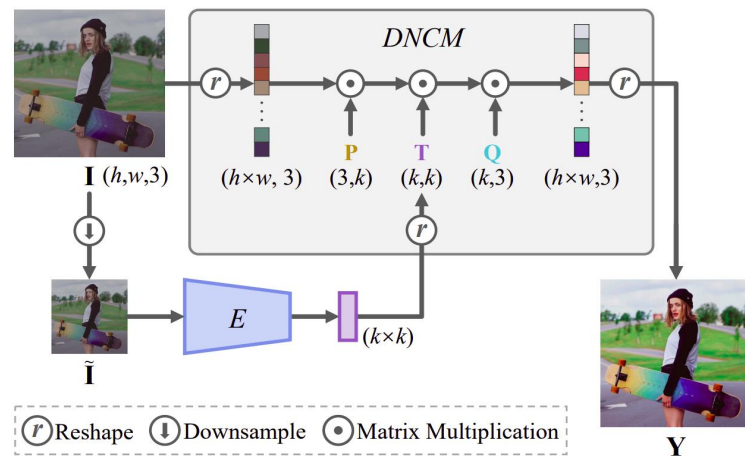




Figure 5. Proposed DNCM Module

Figure 4, 5 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf

How to Train the Models?

- What parameters do we have to train?
- A unified encoder that gets a resized (256x256) image and outputs d_x and r_x to be used in the DNCM modules.
- Two P ($3 \times k$) and two Q ($k \times 3$) matrices for sDNCM and nDNCM.



Training Objective

$$\min \mathbb{E}_{\mathbf{I}_c, \mathbf{I}_s, \mathbf{G}_s \sim p_{\mathbf{I}}} \left[\left| \mathbf{G}_s - S_2(S_1(\mathbf{I}_c), \mathbf{I}_s) \right| \right],$$

\mathbf{I}_c is the content image, \mathbf{I}_s is the style image.
 S_1 is the the color normalization, and S_2 is the stylization.

$$\min \mathbb{E}_{\mathbf{I} \sim p_{\mathbf{I}}} \left[\sum_i^n \sum_j^n \left| \mathbf{I}_j - S_2(S_1(\mathbf{I}_i), \mathbf{I}_j) \right| \right]$$

\mathbf{I}_i is the content image, \mathbf{I}_j is the style image.
 \mathbf{I}_i and \mathbf{I}_j are obtained from the same image with random filters.
Hence, their "content" should be same.
 S_1 is the the color normalization, and S_2 is the stylization.

Training Objective

$$\min \mathbb{E}_{\mathbf{I} \sim p_{\mathbf{I}}} \left[\sum_i^n \sum_j^n \left| \mathbf{I}_j - S_2(S_1(\mathbf{I}_i), \mathbf{I}_j) \right| \right]$$

Given any fixed S_1 , the optimal S_2^* should satisfy:

$$\mathbf{I}_j = S_2^*(S_1(\mathbf{I}_j), \mathbf{I}_j) = S_2^*(S_1(\mathbf{I}_i), \mathbf{I}_j).$$

WARNING!

Using end-to-end CNNs (like most of the autoencoders) for S_1 and S_2 leads to a serious problem for this objective function.

S_2 can easily learn to output the identity of its right-side input, ignoring the output of S_1 , perfectly optimizing the objective with a trivial solution.

So, $S_2(X, I_j) = I_j$ can be obtained with CNNs, which is not what we want.

Self-Supervised Learning

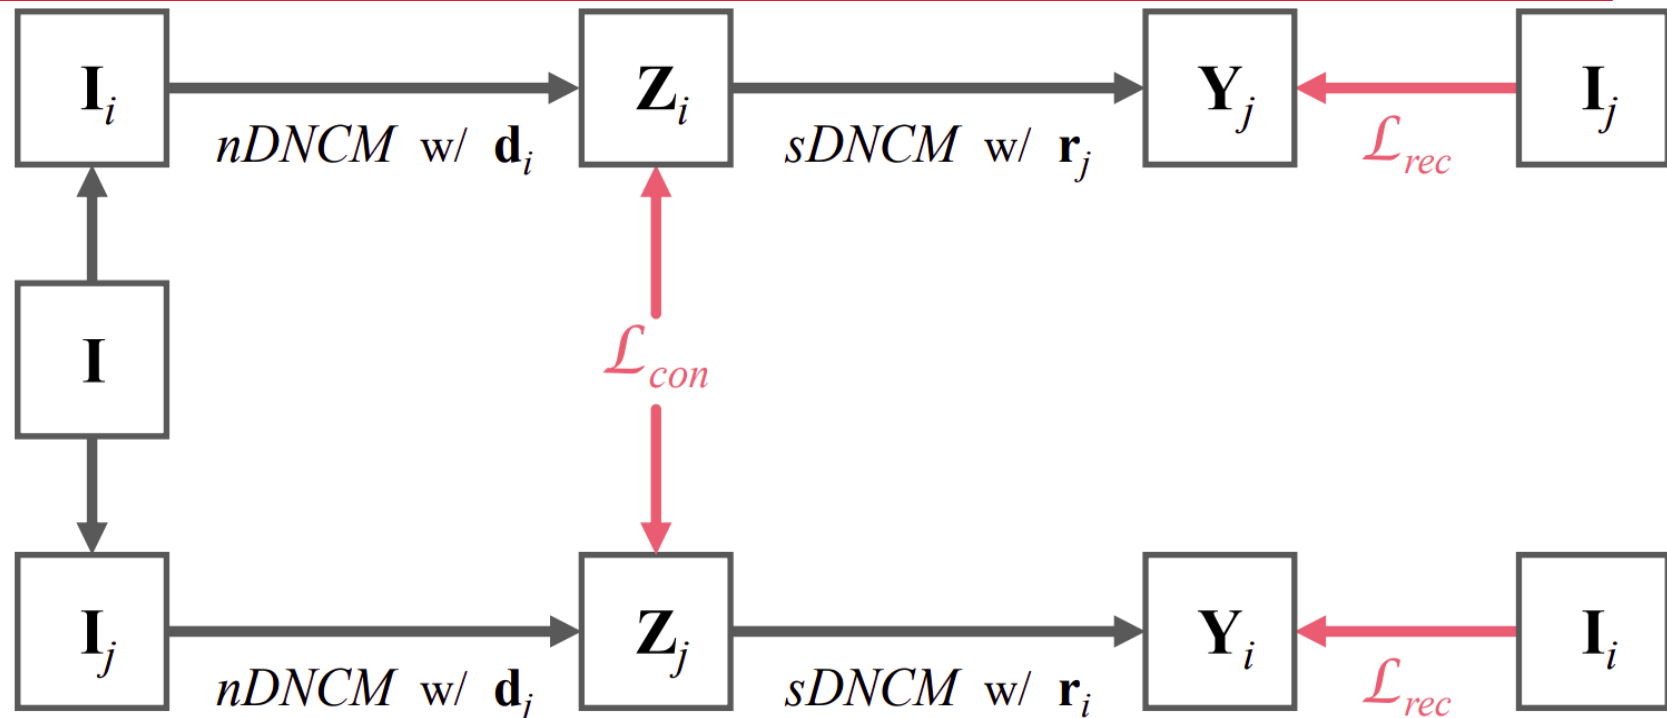


Figure 6. Self-Supervised Learning approach.

Figure 6 (from the presented paper)

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

Self-Supervised Learning

- Obtain I_i and I_j from the same I by applying perturbations (random image filters).
- Get the d_i/r_i and d_j/r_j from the encoder outputs. Apply d_i and d_j to I_i and I_j , respectively, with nDNCM to get normalized color space outputs Z_i and Z_j .
- Apply r_i and r_j to Z_j and Z_i , respectively, to get stylized images.

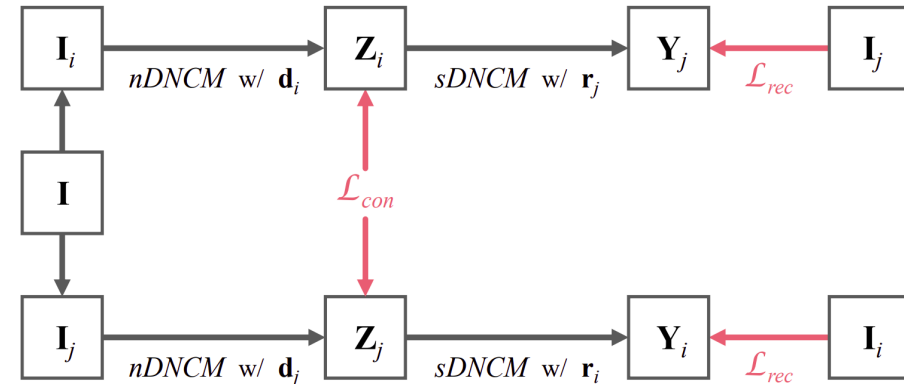


Figure 6. Self-Supervised Learning approach.

Training Objectives

We want Z_i and Z_j to be exactly the same.

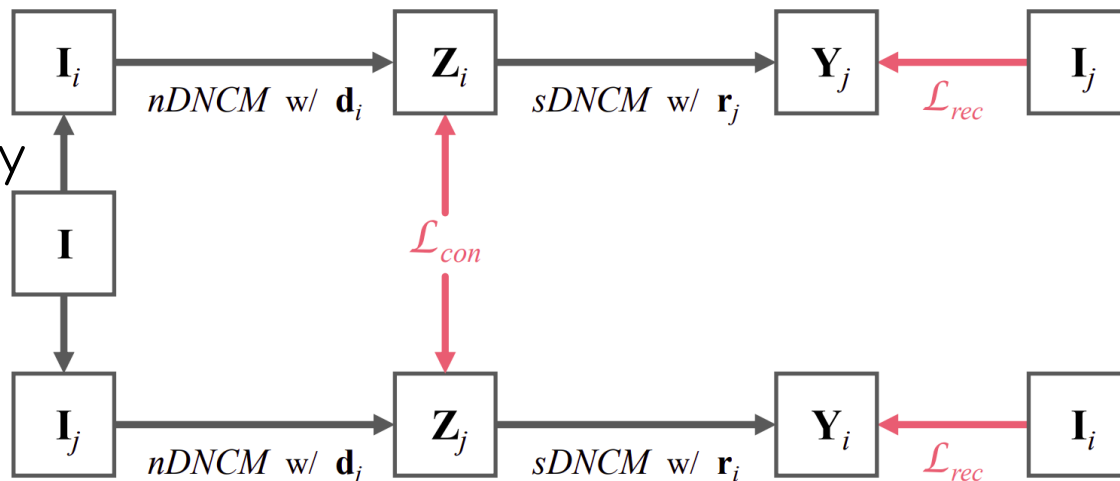


Figure 6. Self-Supervised Learning approach.

We want Y_i and I_j to be exactly the same.

We want Y_j and I_i to be exactly the same.

$$\begin{aligned}\mathcal{L}_{con} &= \|Z_i - Z_j\|_2 \\ &= \|nDNCM(I_i, d_i) - nDNCM(I_j, d_j)\|_2.\end{aligned}$$

$$\mathcal{L}_{rec} = \|Y_i - I_i\|_1 + \|Y_j - I_j\|_1.$$

$$\mathcal{L} = \mathcal{L}_{rec} + \lambda \mathcal{L}_{con},$$

Figure 6 (from the presented paper)

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

4- Experiments



Datasets and Implementation

- MS COCO dataset (*) for training.
- 5000 LUTs, and input color perturbations (**) for getting I_i and I_j .
- EfficientNet-B0 is used with 256x256 input size as the encoder.
- K is picked as 16.
- Trained with Adam for 32 epochs.

* Microsoft COCO: Common Objects in Context
<https://arxiv.org/pdf/1405.0312>

** Harmonizer: Learning to Perform White-Box Image and Video Harmonization
<https://arxiv.org/pdf/2207.01322>

*** EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks
<https://arxiv.org/pdf/1905.11946>

Quantitative Evaluation

- They observed prior metrics does not measure quality well.
- For style similarity, they trained a discriminator on a labeled dataset (with 700+ style classes) to predict the style similarity.
- For content similarity, they compute SSIM on image edges. Prior works used HED (*), but since it is often incorrect and predict only rough edges, they replaced it with LDC (**).

Similarity Comparison

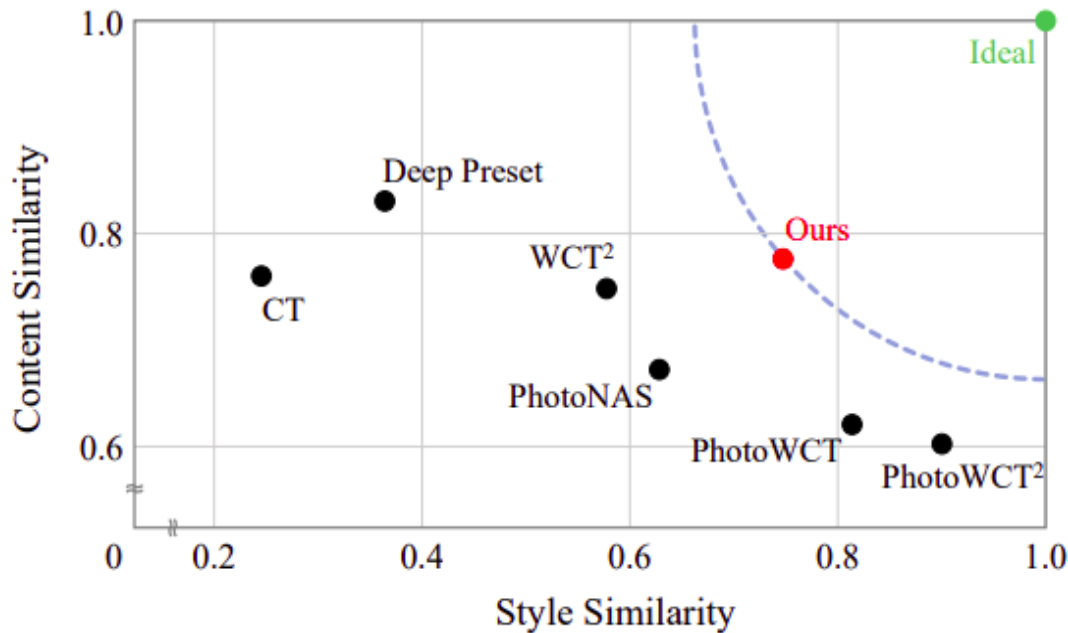


Figure 7. Similarity comparisons with prior methods

Figure 7 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

Similarity Comparison

Method	CT [43]	PhotoWCT [34]	WCT ² [54]	PhotoNAS [1]	PhotoWCT ² [6]	Deep Preset [19]	Ours
Average Ranking ↓	4.97	5.75	2.67	3.39	4.11	5.30	1.81

Figure 8. Average Human Ranking (subjective) Comparisons



Figure 9. Qualitative Comparisons

Figure 8, 9 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

Inference Time Comparison

Method	GPU Inference Time ↓ / Memory ↓				Model Size ↓
	FHD (1920 × 1080)	2K (2560 × 1440)	4K (3840 × 2160)	8K (7680 × 4320)	Number of Parameters
PhotoWCT [34]	0.599 s / 10.00 GB	1.002 s / 16.41 GB	OOM	OOM	8.35 M
WCT ² [54]	0.557 s / 18.75 GB	OOM	OOM	OOM	10.12 M
PhotoNAS [1]	0.580 s / 15.60 GB	0.988 s / 23.87 GB	OOM	OOM	40.24 M
Deep Preset [19]	0.344 s / 8.81 GB	0.459 s / 13.21 GB	1.128 s / 22.68 GB	OOM	267.77 M
PhotoWCT ² [6]	0.291 s / 14.09 GB	0.447 s / 19.75 GB	1.036 s / 23.79 GB	OOM	7.05 M
Ours	0.013 s / 1.96 GB	0.016 s / 1.96 GB	0.019 s / 1.96 GB	0.061 s / 1.96 GB	5.15 M

Figure 10. Inference Time, GPU memory and Model Size comparisons with RTX3090 GPU.

Ablation Studies

k	2	4	8	<u>16</u>	32
Style Similarity \uparrow	0.128	0.510	0.636	<u>0.746</u>	0.769
Content Similarity \uparrow	0.765	0.823	0.781	<u>0.771</u>	0.764

Figure 11. Effect of the "k" hyper-parameter on the similarity values.

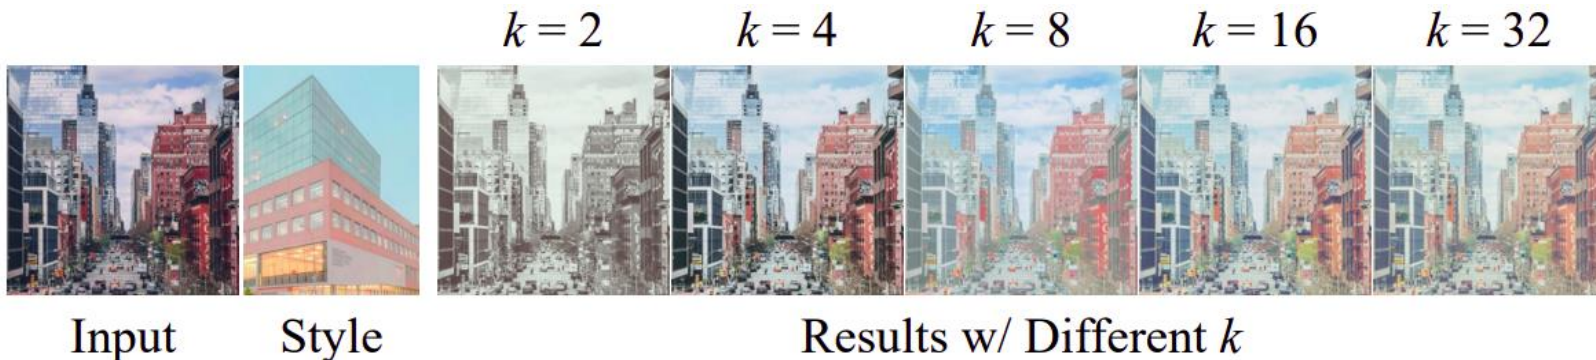


Figure 12. Effect of the "k" hyper-parameter visualized.

Ablation Studies

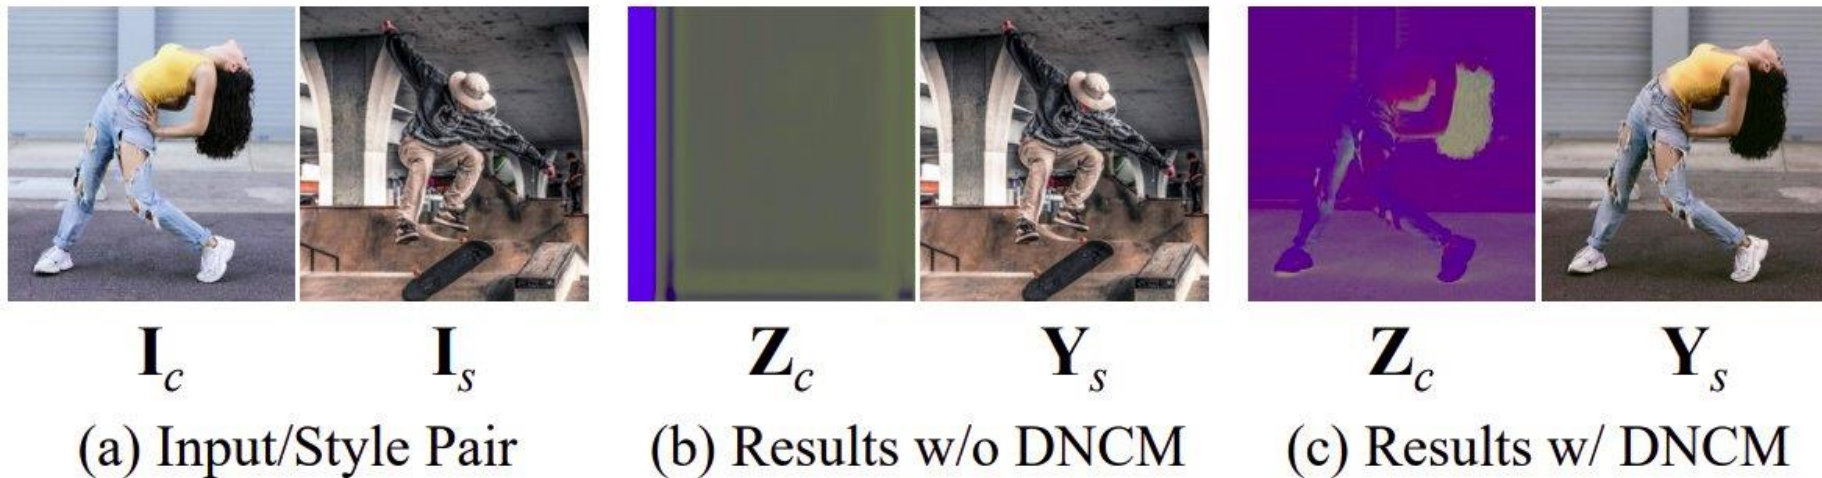


Figure 13. Effect of not using DNCM. Demonstration of the tendency to escaping to a trivial solution.

Ablation Studies




Figure 14. Effect of not using consistency loss

Figure 14 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

5- Main Contribution and Pros/cons



Main Contribution

- DNCM (Deterministic Neural Color Mapping) to avoid artifacts.
- Two-stage approach, enables caching and provides high speed.
- Self-supervised learning approach.
- Custom/improved and better quantitative similarity metrics.



Advantages

- The approach can scale efficiently with high resolution inputs.
- DNCM and two-stage approach significantly reduced the computational expense (Low GPU memory usage and approximately 28x speed up).
- Most of the major artifacts avoided with DNCM. It is crucial for keeping the delicate content such as text readability.
- Without fine-tuning, the model can be used in other tasks such as low-light enhancement.



Disadvantages

- The approach is observed to increase JPEG artifacts.
- If some content colors does not exist on the style image, result may be unwanted.
- Local-adaptive color style transfer is not supported.



Input

Ours

Figure 15. Demonstration of increased JPEG artifacts



Input

Ours

Figure 16. Unsatisfactory outputs



Input

Style

PhotoNAS

PhotoWCT²

Ours

Figure 17. Demonstration of the lack of local-adaptive mapping

Figure 15, 16, 17 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

6- Conclusion



Conclusion

- Using a two-stage approach may tremendously speed up the inference time when switching styles.
- The self-supervised approach can be utilized with DNCM approach, which also avoids some distortions.
- No need for deep and heavy architectures for color style transfer.



Figure 18. Output comparison of the proposed method including the inference times

Figure 18 (from the presented paper):

https://openaccess.thecvf.com/content/CVPR2023/papers/Ke_Neural_Preset_for_Color_Style_Transfer_CVPR_2023_paper.pdf



METU

References (1 of 2)

- Ke, Z., Liu, Y., Zhu, L., Zhao, N., & Lau, R. W. H. (2023, March 23). Neural Preset for Color Style Transfer. arXiv.org. <https://arxiv.org/abs/2303.13511>
- Ke, Z., Sun, C., Zhu, L., Xu, K., & Lau, R. W. H. (2022, July 4). Harmonizer: Learning to Perform White-Box Image and Video Harmonization. arXiv.org. <https://arxiv.org/abs/2207.01322>
- Lin, T. Y., Maire, M., Belongie, S., Bourdev, L., Girshick, R., Hays, J., Perona, P., Ramanan, D., Zitnick, C. L., & Dollár, P. (2014, May 1). Microsoft COCO: Common Objects in Context. arXiv.org. <https://arxiv.org/abs/1405.0312>



References (2 of 2)

- Soria, X. G., Pomboza-Junez, G., & Sappa, N. D. (2022, January 1). LDC: Lightweight Dense CNN for Edge Detection. IEEE Access. <https://doi.org/10.1109/access.2022.3186344>
- Tan, M., & Le, Q. (2019, May 28). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. arXiv.org. <https://arxiv.org/abs/1905.11946>
- Xie, S., & Tu, Z. (2015, December 1). Holistically-Nested Edge Detection. <https://doi.org/10.1109/iccv.2015.164>



Attribution

Keep this page and the following notice AS IS.

This presentation has been prepared with [METU Presentation Template](#) by Devrim Çavuşoğlu licensed under the [CC BY-SA 4.0](#). Also read, [the full LICENSE](#) content.

Refer to github.com/devrimcavusoglu/metu-presentation-template

Thank you for listening
Q & A

